



ARGUS

THE AGENT THAT WATCHES, OWNS, AND DELIVERS

Product Requirements Document

v0.1 • May 2026

AUTHOR

Nikhil Kumar Gupta

STATUS

Draft for review

TARGET LAUNCH

Private beta — Q4 2026

POSITIONING

Hermes killer

The next paradigm is outcomes, not tasks.

Argus is an autonomous agent network that owns business outcomes end-to-end. Unlike Hermes Agent, which learns inside a single user's closed loop, Argus composes specialist agents that learn across users via federated skill graphs, act on their own initiative, and prove every action via cryptographic lineage. It is positioned for teams and regulated work — the segment Hermes structurally cannot serve.

The bet: the unit of value in AI agents shifts from *tokens generated* to *outcomes delivered*. Argus prices on that unit. The first agent to make outcome-SLA pricing default wins the category.

NORTH STAR

\$ARR/agent

Revenue produced per
active agent / month

ACTIVATION

T-72h

Time to first owned outcome

TRUST

>99.9%

Signed actions w/ valid
provenance

Why now, and why this shape.

The disruption pattern

Each market-leading agent has won on a new axis, not by beating the prior leader on its own terms. Cursor reframed AI from chat to in-the-IDE (proximity). OpenClaw reframed it from tool to ecosystem (breadth). Hermes reframed it from executor to learner (depth). The next leader will not learn faster — it will **act on its own, learn from everyone, and prove what it did.**

What Hermes structurally cannot solve

Hermes’ closed per-user loop is its identity. It cannot become federated without abandoning the privacy story that sold the architecture. It cannot become proactive without breaking its session contract. It cannot become verifiable without giving up the silent self-rewriting that defines “the agent that grows with you.” Each of these is a one-way door for Hermes. That asymmetry is the opening.

Market timing

Three conditions just lined up. **(1) Inference economics:** Haiku-class and GPT-5-mini-class models make always-on agents viable at SaaS unit economics. **(2) Buyer fatigue:** the OpenClaw CVE wave (9 CVEs in 4 days, March 2026) and the ClawHub 12% malware audit have created enterprise demand for a verifiable alternative. **(3) Standards:** MCP has matured to where outcome contracts can be expressed portably across tools, removing the lock-in objection.

Competitive landscape

Dimension	OpenClaw	Hermes Agent	Argus
Core thesis	Ecosystem breadth	Self-improvement	Outcome ownership
Learning unit	Per-skill (manual)	Per-user (closed loop)	Federated, cross-user
Initiative	Reactive	Reactive	Proactive (event-driven)
Pricing model	Open source / seats	Tokens	Outcome SLA
Trust model	Community signals	Implicit (drift risk)	Signed actions + content-addressed skills
Failure mode	Supply-chain malware	Silent regression	Reversion via signed lineage
Best for	Tinkerers, indie devs	Power users, solo operators	Teams + regulated work

Who hurts, and how much.

Problem statement

Operators (founders, agency owners, solo PMs, RevOps leads) buy AI agents to free up calendar time and ship outcomes, but every current agent — including Hermes — operates at the task layer. The user still owns the goal, the orchestration, the recovery from failure, and the audit trail. The agent saves minutes, not weeks. For regulated buyers the gap is worse: they cannot deploy a self-modifying agent into a SOX, HIPAA, or PCI environment because no current agent can prove what it did or why.

Primary persona — “Maya, the operator”

Solo or small-team operator running 3–5 workflows weekly (outreach, content, reporting, lead qualification). Tool-rich, time-poor. Currently stitching Hermes + 6 MCPs + a Notion doc. Will pay \$500–\$2,000/mo for an agent that *owns a number* instead of helping with steps.

Secondary persona — “Rahul, the head of RevOps”

Mid-market RevOps lead at a 50–500-person SaaS. Needs to deploy agents into a regulated stack with audit, RBAC, and rollback. Has rejected OpenClaw post-CVE and refuses to deploy Hermes because security cannot certify a self-modifying runtime. Buyer for the enterprise tier.

An agent shaped like an employee, not a tool.

“Argus is the hundred-eyed watcher of myth — the agent that never sleeps, learns from every eye, and signs every glance.” Operationally: you hand Argus an outcome and a budget, and it reports back when the outcome lands or when it genuinely needs you. You stop opening the app. You start reviewing weekly summaries.

Product principles

Outcomes > tasks	Every contract is expressed as a measurable outcome with success criteria, budget, and deadline. Tasks are an implementation detail Argus owns.
Proactive by default	Argus initiates. Notifications go from agent to user, not the other way around. The default UI is the inbox digest, not a chat window.
Federated, never centralized	Skill graphs train across opt-in users via differential privacy on embeddings. No raw data leaves the user’s tenant.
Signed or it didn’t happen	Every action Argus takes is cryptographically signed and content-addressed. Replay, diff, and revert are first-class.
Composable specialists	Argus is a fleet of small specialist agents — outreach, reporting, code-review, finance-recon — that compose at runtime. No monolithic “Argus model.”

What ships in v1.

Outcome Contracts	Declarative spec for what Argus must deliver: success metric, deadline, budget, escalation policy. Renders as a one-page artifact the user signs. The contract is the source of truth — agents read it, not chat history.
Specialist Fleet	Pre-trained micro-agents per outcome family: Outbound, Inbound Qualification, Weekly Reporting, Code Review, Reconciliation, Content Repurposing. Each agent has a public skill graph; new ones can be forked and published.
Federated Skill Graph	Cross-user pattern learning over differentially-private skill embeddings. Result: a brand-new Argus deployment ships with the equivalent of a 6-month-experienced operator on day one.
Lineage Ledger	Append-only signed log of every action Argus takes — every tool call, every model invocation, every decision branch. Inspectable, diff-able, revertable. The audit story for regulated buyers.
Initiative Engine	Event-driven runtime that triggers agents on calendar, inbox, repo, and CRM changes. The user is interrupted only when the contract explicitly requires human input.
Outcome-SLA Pricing	Default plan is \$X per delivered outcome with a monthly floor. Customers see exactly what they paid for — and refunds are automatic when Argus fails to meet the SLA.

Four layers, one contract.

Contract Layer	Outcome spec, success criteria, budget, escalation rules. Source of truth. Stored per-tenant; versioned with semantic diff.
Fleet Layer	Pool of specialist agents. Each agent is a content-addressed bundle (model + tool list + skill graph + policy). Composed at runtime against the contract.
Lineage Layer	Signed, append-only event log of every action. Ed25519 signing keys per tenant. Public verification endpoint. Lineage is the artifact auditors consume.
Federation Layer	Off-tenant differential-privacy training over skill embeddings (not raw events). Privacy budget tracked per user. Opt-in only; opt-out is one toggle and removes the user's past contributions.

Success metrics

North Star	\$ARR generated per active agent per month
Activation	Time to first delivered outcome < 72 hours from signup
Retention	> 95% logo retention at 90 days, > 85% at 12 months
Trust	> 99.9% of actions land in the lineage ledger with valid signatures; zero post-mortem-grade trust
Federation lift	> 20% reduction in time-to-first-outcome for new users with federation opted in vs. cold start
Outcome SLA	> 92% of contracts meet stated success criteria within deadline

What can kill this.

Risk	Likelihood	Impact	Mitigation
Federated learning leaks private data via skill graphs	Medium	Critical	Differential privacy on skill embeddings; opt-in contribution; red-team budget at ev
Outcome-SLA pricing under-prices long-tail workflow	High	Medium	Hybrid: outcome floor + seat ceiling; offer downgrade to seat-only after 90 days
Enterprise procurement cycle stalls Q4 launch	High	High	Design-partner program (5 customers, no procurement) running in parallel; SOC 2
Hermes ships federated learning workflow	Medium	High	Lead on outcome ownership + verifiability; published lineage spec creates standar
Skill-marketplace malware repeats OpenClaw failure	Medium	Critical	Content-addressed only; no string-name resolution; signed publisher identities; sa

Go-to-market

Three concentric circles. **Inner:** 5 design partners across RevOps, agency, and fintech-back-office — hand-onboarded, contract-priced, used as case studies. **Middle:** a public skill marketplace where operators can fork specialist agents (this is the OpenClaw playbook, executed with content-addressed trust from day one). **Outer:** developer mindshare via an open lineage spec — make it the “OpenTelemetry of agent actions” so anyone can verify any agent, not just Argus.

Distribution wedge: regulated verticals where Hermes is structurally unsellable. Finance ops, healthcare intake, legal review, SOX testing. Argus enters where Hermes is banned and expands sideways from there.

Roadmap

Phase	Window	Goal	Exit criteria
P0 — Foundations	Jun–Aug 2026	Outcome contract spec + signed-action workflow	10 contracts executed end-to-end with full lineage
P1 — Design partners	Sep–Nov 2026	5 paying design partners, 3 verticals	\$50K MRR, NPS > 40, < 1 critical security incident
P2 — Federated v1	Dec 2026 – Feb 2027	Skill graphs trained across opt-in users	20% reduction in time-to-first-outcome vs. cold start
P3 — General availability	Mar 2027	Self-serve onboarding + SOC 2 Type II	100 paying customers, \$250K MRR, churn < 3%

In the myth, Hermes lulled Argus’ hundred eyes to sleep and killed him. This time, Argus stays awake.